

The Automation Playbook Workbook

One workbook that becomes your one-page Automation Playbook — the unit is always the **task**, never the tool. Fill it as the course goes; by the capstone the six tabs **are** your Playbook. The gray rows are the worked example — the 45-person firm's invoice-coding task — write your own underneath.

The automation DRAFTS. A named human signs before anything is posted, paid, sent, or filed — you never automate the signature. Anything that touches money, compliance, or an action you can't take back stays a draft a person signs, every time, whatever the score.

1. **Score** your candidate tasks on five plain questions and read the verdict — Module 1.
2. **Choose** the lightest rung of the ladder that does the job — Module 2.
3. **Author** the SOP your automation runs on, and point the SOP-Execution Agent at it — Module 3.
4. **Launch** the smallest safe version, with the human's sign-off designed in — Module 4.
5. **Evaluate** it: test the exceptions, arm the pivot switch, measure baseline vs after — Module 5.

S · SCORE

Automatability Scorecard

C · CHOOSE

Approach Pick

A · AUTHOR

Task SOP + Agent

L · LAUNCH

Launch Card

E · EVALUATE

Test / Rollback / Yield

Your Starting Snapshot — the task you’re bringing

You arrive from Tasks, Not Jobs with one task you already decided to automate — your First Move.

Field	Your First Move	Example — the firm's invoice-coding task
The task (verb + noun)		<i>Code incoming vendor invoices to the right GL account</i>
Who does it today		<i>The AP clerk, in Finance — about 6 hours a week</i>
Who approves its output today (the signer)		<i>The controller — signs the coding before it posts</i>
What the task touches (data class)		<i>Payroll + vendor terms → money · Confidential</i>
Why you chose it as your First Move		<i>Execution work, high time-drain, decided 'automate' in Tasks, Not Jobs</i>

The Automatability Scorecard

First three columns: higher is more automatable. Last three (weird / judgment / systems): higher pushes AWAY from automation. Verdict bands are Build now / Not yet / Keep it human — never a traffic light. The band is a floor, not a ceiling.

[illegible]

The Approach Pick

Four rungs, lightest first: 1 Prompt/agent recipe · 2 No-code workflow · 3 Connected integration · 4 Custom build. Start on the lowest rung that does the job.

Task	Rung (start lowest that works)	Why this rung	What a later climb would take (the trigger)
Code vendor invoices to GL accounts	1 — Prompt / agent recipe	Structured, repeatable, feeds one system; the SOP-Execution Agent drafts, the controller signs. A money-touch does NOT push it up.	Climb to a 2 — No-code workflow only if volume roughly triples, a second entity needs it, or manual copy-paste becomes the bottleneck.

Your Task SOP

You can't automate what you can't write down. Point the SOP-Execution Agent at this; it drafts from these steps, a named human signs. The exceptions become your test cases in Module 5.

SOP part	Your task	Example — invoice coding
Task name		<i>Code incoming vendor invoices to the right GL account</i>
Trigger (what starts it)		<i>A new vendor invoice lands in the AP inbox</i>
Inputs (what the agent reads)		<i>The invoice; the vendor→category list; the chart of accounts</i>
Steps (numbered)		<i>1) Read vendor + line items. 2) Match vendor to its usual category. 3) Apply the GL rule. 4) Produce a draft coding with the amount and account.</i>
Exception rules (if unsure, FLAG — don't guess)		<i>Split invoice → FLAG. New vendor not on the list → FLAG. Credit / credit memo → FLAG. Amount over threshold → FLAG.</i>
Sign-off (the named human, before anything posts)		<i>The controller reviews the draft, edits if needed, and posts. The agent never posts, pays, or touches the ledger.</i>

The Launch Card

If a cell is blank, the automation doesn't go live. The approval step is designed in from day one, never bolted on.

Launch Card field	Your automation	Example — invoice coding
Scope (as small as still works)		<i>One vendor category only (software subscriptions)</i>
Human approval step (who signs, before what)		<i>The controller signs every draft BEFORE it posts — no auto-post, ever</i>
Boundary (where it runs)		<i>Company accounts + approved AI; drafts only — never the ledger or payment</i>
Data class (what it touches)		<i>Confidential — payroll + vendor terms</i>
Go-live plan (one task, one team, one week)		<i>Week 1, AP + controller only; the rest of the team is told what it is and isn't</i>
Fallback if it misbehaves (the pivot switch)		<i>Turn the agent off; the AP clerk codes manually — the way it always worked</i>

Test · Rollback · Yield

Test the exceptions your scorecard flagged. Every automation gets a pivot switch. The yield is a labeled estimate until the 90-day checkpoint proves it — hand the honest read to Signal vs Noise.

Field	Your automation	Example — invoice coding (numbers are a labeled synthetic example)
Test cases (MUST include the flagged exceptions)		<i>Split invoice · new vendor · credit memo · over-threshold — plus ordinary ones</i>
Pivot switch (retune → redirect → stop) + who flips it		<i>Controller flips it: retune the SOP, redirect odd cases to manual, or stop and code by hand</i>
Manual fallback (the way it always worked)		<i>The AP clerk codes invoices by hand — no automation ships without this</i>
Baseline — counted BEFORE (label the number)		<i>~6 hrs/week on the pilot category (measured over 2 weeks) — synthetic</i>
After — measured at the checkpoint		<i>~2 hrs/week drafting + review — synthetic</i>
Yield (label it an ESTIMATE until day 90)		<i>~4 hrs/week back on this one category (estimate) — no 'average', no benchmark</i>
90-day checkpoint date (on the calendar)		<i>Go-live + 90, on the controller's calendar, baseline pasted into the invite</i>
What this yield funds next (the SCALE loop)		<i>Automation #2 — the proven hours pay for the next task; no big-bang spend</i>